

Message Passing Paradigm
MPI : Distributed memory systems

MPI stands for **Message Passing Interface**.

- Library of subroutines/functions - not a language.
- Programmer insert appropriate MPI subroutine/function calls, compile and finally link with MPI message passing library.

It is possible to write fully-functional message-passing programs by using **only six routines**. The first three examples.

Typical MPI Program Structure

MPI include file (header)

Initialize MPI environment

Message Passing Calls

Terminate MPI Environment

All MPI routines, data-types, and constants are prefixed by “MPI_”.

The return code for successful completion is MPI_SUCCESS.

Six MPI Functions

These six functions allows us to write many programs:

MPI_Init()

MPI_Finalize()

MPI_Comm_size()

MPI_Comm_rank()

MPI_Send()

MPI_Recv()

MPI Processes

- MPI is process oriented - program consists of multiple processes, each corresponding to one processor.
- In practice - runs its own copy of the same code (SPMD).
- MPI process and threads - MPI process can contain a single thread or multiple threads.
- MPI processes are identified by their **ranks**
 - *If total $nprocs$ processes in computation, rank ranges from 0, 1, ..., $nprocs-1$.*
 - *$nprocs$ does not change during computation.*

Example

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char **argv)
{
    int my_rank, nprocs;

    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);
    MPI_Comm_size(MPI_COMM_WORLD, &nprocs);
    printf("This is process %d among %d processes\n",
           my_rank, nprocs);
    MPI_Finalize();
    return 0;
}
```

MPI Environment

- MPI_Init
- MPI_Comm_size
- MPI_Comm_rank
- MPI_Finalize

- MPI_Wtime

Returns an elapsed wall clock time in seconds on the calling processor

Initialization and Termination

- `MPI_Init()` initializes MPI environment
- Must be called before any other MPI routine
- Called only once; subsequent calls are erroneous.

- `MPI_Finalize()` cleans up MPI environment
 - *Must be called before exits.*
 - *No other MPI routine can be called after this call*
 - *Exception: `MPI_Initialized()`; `MPI_Get_version()`, `MPI_Finalized()`.*

- Abnormal termination: `MPI_Abort()`
 - *Makes a best attempt to abort all tasks*

Communicators and Process Groups

- A communicator defines a *communication domain* - a set of processes that are allowed to communicate with each other.
- Information about communication domains is stored in variables of type `MPI_Comm`.
- Communicators are used as arguments to all message transfer MPI routines.
- A process can belong to many different communication domains.
- MPI defines a default communicator called `MPI_COMM_WORLD` which includes all the processes.
- `MPI_Comm_size`
Determines the number of processes in the group associated with a communicator
- `MPI_Comm_rank`
Determines the rank of the calling process within the communicator

Communicators

- How many CPUs: `MPI_COMM_SIZE()`
- Who am I: `MPI_COMM_RANK()`
- Can be used for data decomposition etc.
 - *Suppose we know total number of grid points, total number of cpus and current cpu id.*
 - *We can calculate which portion of data current cpu is to work on.*
- Ranks used to specify source and destination of communications.

my_rank value different on different processors !

Second Example

```
#include <mpi.h>
#include <stdio.h>
#include <string.h>

int main(int argc, char **argv)
{
    char message[256];
    int my_rank;
    MPI_Status status;

    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);
    if(my_rank==0) {
        strcpy(message, "Hello");
        MPI_Send(message, strlen(message)+1, MPI_CHAR, 1, 99, MPI_COMM_WORLD);
    }
    else if(my_rank==1) {
        MPI_Recv(message, 256, MPI_CHAR, 0, 99, MPI_COMM_WORLD, &status);
        printf("Process %d received: %s\n", my_rank, message);
    }
    MPI_Finalize();
    return 0;
}
```

6 MPI functions:

MPI_Init()

MPI_Finalize()

MPI_Comm_rank()

MPI_Comm_size()

MPI_Send()

MPI_Recv()

➤ Point-to-point communications

- *Involves a sender and a receiver, one processor to another processor*
- *Only the two processors participate in communication*

➤ Collective communications

- *All processors within a communicator participate in communication
e.g. Barrier, reduction operations, gather, ...*

```
...  
MPI_Send(message, strlen(message)+1, MPI_CHAR,  
1, 99, MPI_COMM_WORLD);  
MPI_Recv(message, 256, MPI_CHAR, 0, 99, MPI_COMM_WORLD, &status);  
...
```

- Message data: what to send, what to receive?
 - *Where is the message? Where to put it?*
 - *What kind of data is it? What is the size?*
- Message envelope: where to send/receive?
 - *Sender, receiver*
 - *Communication context*
 - *Message tag.*

Send

```
int MPI_Send(void *buf,int count,MPI_Datatype datatype,  
             int dest, int tag, MPI_Comm comm)  
MPI_SEND (BUF,COUNT,DATATYPE,DEST,TAG,COMM,IERROR)
```

- `buf` – memory address of start of message
- `count` – number of data items
- `datatype` – what type each data item is (integer, character, double, float ...)
- `dest` – rank of receiving process
- `tag` – additional identification of message
- `comm` – communicator, usually `MPI_COMM_WORLD`

```
char message[256];  
MPI_Send(message,strlen(message)+1,MPI_CHAR,  
1,99,MPI_COMM_WORLD);
```

Receive

```
int MPI_Recv(void *buf,int count,MPI_Datatype datatype,int source,int tag,  
MPI_Comm comm,MPI_Status *status)
```

```
MPI_RECV(BUF,COUNT,DATATYPE,SOURCE,TAG,COMM,STATUS,IERROR)
```

- `buf` – initial address of receive buffer
- `count` – number of elements in receive buffer (size of receive buffer)
- `datatype` – data type in receive buffer
- `source` – rank of sending process
- `tag` – additional identification for message
- `comm` – communicator, usually `MPI_COMM_WORLD`
- `status` – object containing additional info of received message
- `ierror` – return code

```
char message[256];  
MPI_Recv(message,256,MPI_CHAR,0,99,MPI_COMM_WORLD,&status);
```

Actual number of data items received can be queried from status object; it may be smaller than count, but cannot be larger (if larger → overflow error).

Basic MPI Data Types

MPI datatype	C datatype
MPI_CHAR	signed char
MPI_SHORT	signed short
MPI_INT	signed int
MPI_LONG	signed long
MPI_UNSIGNED_CHAR	unsigned char
MPI_UNSIGNED_SHORT	unsigned short
MPI_UNSIGNED	unsigned int
MPI_UNSIGNED_LONG	unsigned long int
MPI_DOUBLE	double
MPI_FLOAT	float
MPI_LONG_DOUBLE	long double
MPI_BYTE	
MPI_PACKED	

Essentials of communication


- Sender must specify valid destination.
- Sender and receiver data type, tag, communicator must match.
- Receiver returns extra (status) parameter to report info regarding message received.
- Sender specifies size of sendbuf; receiver specifies *upper bound* of recvbuf.

Message Passing

- Both standard send and receive functions are *blocking*
- `MPI_Recv` returns only after receive buffer contains requested message
- `MPI_Send` may or may not block until message received (usually blocks)
- May lead to deadlock

MPI :
Integration using trapezoidal rule

Parallelizing – Integration using Trapezoidal Rule



- *Function f over Problem interval $[A,B]$*
- *Create and Distribute chunks of workload. Each workload consists of its own subinterval of $[a,b]$ which is assigned to each process.*
- *Calculate f for each subinterval*
- *Finally add the f calculated for all the sub intervals to produce result for the complete problem $[A,B]$*

➤ Issues to consider

- *Load balancing - Number of trapezoids (n) are equally divisible across (p).*

➤ Each process needs the following information

- *Rank of the process*
 - *Ability to derive the workload per processor as a function of rank*
- Finally any process can do the summation*

Parallelizing Trapezoidal Rule (SPMD)

➤ Algorithm

Assuming Number of trapezoids n is evenly divisible across p processors

- Calculate:

$$h = \frac{(b - a)}{n}$$

- **Each process calculates its own workload**

- local number of trapezoids (local_n) = n/p
- local starting point (local_a) = $a + (\text{process_rank} * \text{local_n} * h)$
- local ending point (local_b) = $(\text{local_a} + \text{local_n} * h)$

- **Each process calculates its own integral for the local intervals**

- For each of the local_n trapezoids calculate area and sum the area for local_n trapezoids

- **If $\text{PROCESS_RANK} == 0$**

- Receive messages (sub-interval area sum) from all processors
- ADD all sub-interval areas

- **If $\text{PROCESS_RANK} > 0$**

- Send sub-interval area to $\text{PROCESS_RANK}(0)$

.

Basic MPI Calls

- The 6 main MPI calls:
 - *MPI_Init*
 - *MPI_Finalize*
 - *MPI_Comm_size*
 - *MPI_Comm_rank*
 - *MPI_Send*
 - *MPI_Recv*
- Include MPI Header file
 - *#include "mpi.h"*
- Basic MPI Datatypes
 - *MPI_INT, MPI_FLOAT,*

Parallel Trapezoidal Rule

```
#include <stdio.h>
#include "mpi.h"

main(int argc, char** argv) {
    int    my_rank; /* My process rank */
    int    p;      /* The number of processes */
    float  a = 0.0; /* Left endpoint */
    float  b = 1.0; /* Right endpoint */
    int    n = 1024; /* Number of trapezoids */
    float  h;      /* Trapezoid base length */
    float  local_a; /* Left endpoint my process */
    float  local_b; /* Right endpoint my process */
    int    local_n; /* Number of trapezoids for my calculation */
    /*
    float  integral; /* Integral over my interval */
    float  total;    /* Total integral */
    int    source;   /* Process sending integral */
    int    dest = 0; /* All messages go to 0 */
    int    tag = 0;
    MPI_Status status;
```

Trapezoidal Example Adapted from Parallel Programming in MPI P.Pacheco

Parallel Trapezoidal Rule

```
float Trap(float local_a, float local_b, int local_n, float h);    /* Calculate local
integral */
```

```
/* to start up MPI */
```

```
MPI_Init(&argc, &argv);
```

```
/* Get my process rank */
```

```
MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);
```

```
/* Find out how many processes are being used */
```

```
MPI_Comm_size(MPI_COMM_WORLD, &p);
```

```
h = (b-a)/n;    /* h is the same for all processes */
```

```
local_n = n/p; /* So is the number of trapezoids */
```

```
/* Length of each process' interval of * integration = local_n*h. So my interval
* starts at: */
```

```
local_a = a + my_rank*local_n*h;
```

```
local_b = local_a + local_n*h;
```

```
integral = Trap(local_a, local_b, local_n, h);
```

Trapezoidal Example Adapted from Parallel Programming in MPI P.Pach

Parallel Trapezoidal Rule

```
/* Add up the integrals calculated by each process */
if (my_rank == 0) {
    total = integral;
    for (source = 1; source < p; source++) {
        MPI_Recv(&integral, 1, MPI_FLOAT, source, tag, MPI_COMM_WORLD,
&status);
        total = total + integral;
    }
} else {
    MPI_Send(&integral, 1, MPI_FLOAT, dest, tag, MPI_COMM_WORLD);
}
/* Print the result */
if (my_rank == 0) {
    printf("With n = %d trapezoids, our estimate\n", n);
    printf("of the integral from %f to %f = %f\n",
        a, b, total);
}
/* Shut down MPI */
MPI_Finalize();
} /* main */
```

Trapezoidal Example Adapted from Parallel Programming in MPI P.Pacheco

MPI Summary

- MPI Programs are made up of communicating processes
- Each process has its own address space containing its own attributes such as rank, size etc.
- Default communicator is MPI_COMM_WORLD
 - All processes are its members
 - It has a size (the number of processes)
 - Each process has a rank within it