

Scan operation and parallelization

Reduction - recall

2

- we have seen that a sum reduction takes $O(n)$ work to run to completion.
- The parallel sum reduction takes the same order of work (operations) to complete as the serial version.
- Therefore, we say a sum reduction is work-efficient.

Scan (Prefix Sum)

3

- Given an array $A = [a_0, a_1, \dots, a_{n-1}]$ and a binary associative operator \oplus with identity I ,

$$\text{scan}(A) = [I, a_0, (a_0 \oplus a_1), \dots, (a_0 \oplus a_1 \oplus \dots \oplus a_{n-2})]$$

Computing Cumulative Sum

- Prefix sum: if \oplus is addition, then scan on the series

3	1	7	0	4	1	6	3
---	---	---	---	---	---	---	---

returns the series

?????

Scan (Parallel Prefix Sum)

4

- Given an array $A = [a_0, a_1, \dots, a_{n-1}]$ and a binary associative operator \oplus with identity I ,

$$\text{scan}(A) = [I, a_0, (a_0 \oplus a_1), \dots, (a_0 \oplus a_1 \oplus \dots \oplus a_{n-2})]$$

- Prefix sum: if \oplus is addition, then scan on the series

3	1	7	0	4	1	6	3
---	---	---	---	---	---	---	---

returns the series

0	3	4	11	11	15	16	22
---	---	---	----	----	----	----	----

Scan is a simple & useful parallel building block for many parallel algorithms. Cutting example.

Difficult to parallelize !!! Each element depends on the previous.

Scan : types

5

Sum of all the elements that preceded **current element** +

Inclusive Scan

3	1	7	0	4	1	6	3
---	---	---	---	---	---	---	---

returns the series

3	4	11	11	15	16	22	25
---	---	----	----	----	----	----	----

Exclusive Scan

3	1	7	0	4	1	6	3
---	---	---	---	---	---	---	---

returns the series

0	3	4	11	11	15	16	22
---	---	---	----	----	----	----	----

Difference
between
reduction and
scan?

Scan (prefix sum) : important concepts

6

- Input \rightarrow array of elements
- Binary operator
- Identity element/ value \rightarrow when operate with any element gives us back the same element.
- Associative and commutative operator.
- Running Sum of Inputs

Uses for Scan

7

- Sorting
- Lexical Analysis
- String Comparison
- Polynomial Evaluation
- Stream Compaction
- Building Histograms and Data Structures


Serial Scan

```
int input[8] = {3, 1, 7, 0, 4, 1, 6, 3};  
int result[8];  
int running_sum = 0;  
for(int i = 0; i < 8; ++i)  
{  
    result[i] = running_sum;  
    running_sum += input[i];  
}
```

```
prefixSum[0] = 0;  
for(i = 1; i < N; i++)  
    prefixSum[i]=prefixSum[i-1]+arr[j-1];  
  
// result = {0, 3, 4, 11, 11, 15, 16, 22}
```


Reduction


9



Work Complexity → work efficient
Step Complexity → Step efficient

Parallel Reduction

10



Parallel Reduction Complexity

11

- $\text{Log}(N)$ parallel steps, each step S does $N/2^S$ independent ops
 - *Step Complexity is $O(\log N)$*
- For $N=2^D$, performs $\sum_{S \in [1..D]} 2^{D-S} = N-1$ operations
 - *Work Complexity is $O(N)$ – It is *work-efficient**
 - *i.e. does not perform more operations than a sequential algorithm*

Parallel Inclusive Scan (Hillis and Steele)

A Scan Algorithm


13

3	1	7	0	4	1	6	3
---	---	---	---	---	---	---	---

Assume array has shared scope

A Scan Algorithm – Preview

14




Iteration 0, $n-1$ threads

Each \oplus corresponds to a single thread.

Iterate $\log(n)$ times. Each thread adds value *stride* elements away to its own value

A Scan Algorithm – Preview

15




Iteration 1, $n-2$ threads

Each \oplus corresponds to a single thread.

Iterate $\log(n)$ times. Each thread adds value *offset* elements away to its own value

A Scan Algorithm – Preview

16



Iteration i , $n-2^i$ threads

Each \oplus corresponds to a single thread.

Iterate $\log(n)$ times. Each thread adds value *offset* elements away to its own value.

Note that this algorithm operates in-place: no need for double buffering

A Scan Algorithm – Preview

17

3	4	11	11	15	16	22	25
---	---	----	----	----	----	----	----

- We have an **inclusive** scan result
- End of iteration N , array will contain the sum of 2^N input elements at and before the location.
- Input $X[0] \rightarrow$ output $Y[0]$
- First iteration \rightarrow Each position except $Y[0]$ receives the sum of current content and its left element.
- Second iteration \rightarrow Except $Y[0], Y[1]$, receive the sum of its current content and content 2 elements away.


Scan - OpenMP Parallelism

```
for(i = 1; i <= log_of_N; i++)
{
    two_i_1 = 1 << (i-1);
    two_i = 1 << i;
    out = 1 - out;
    in = 1 - out;
    #pragma omp parallel for private(j)
    shared(scanSum, in, out)
    for(j = 0; j < N; j++)
    {
        if(j >= two_i_1)
            scanSum[out][j] = scanSum[in][j] +
scanSum[in][j - two_i_1];
        else
            scanSum[out][j] = scanSum[in][j];
    }
}
```

Work Efficiency Considerations

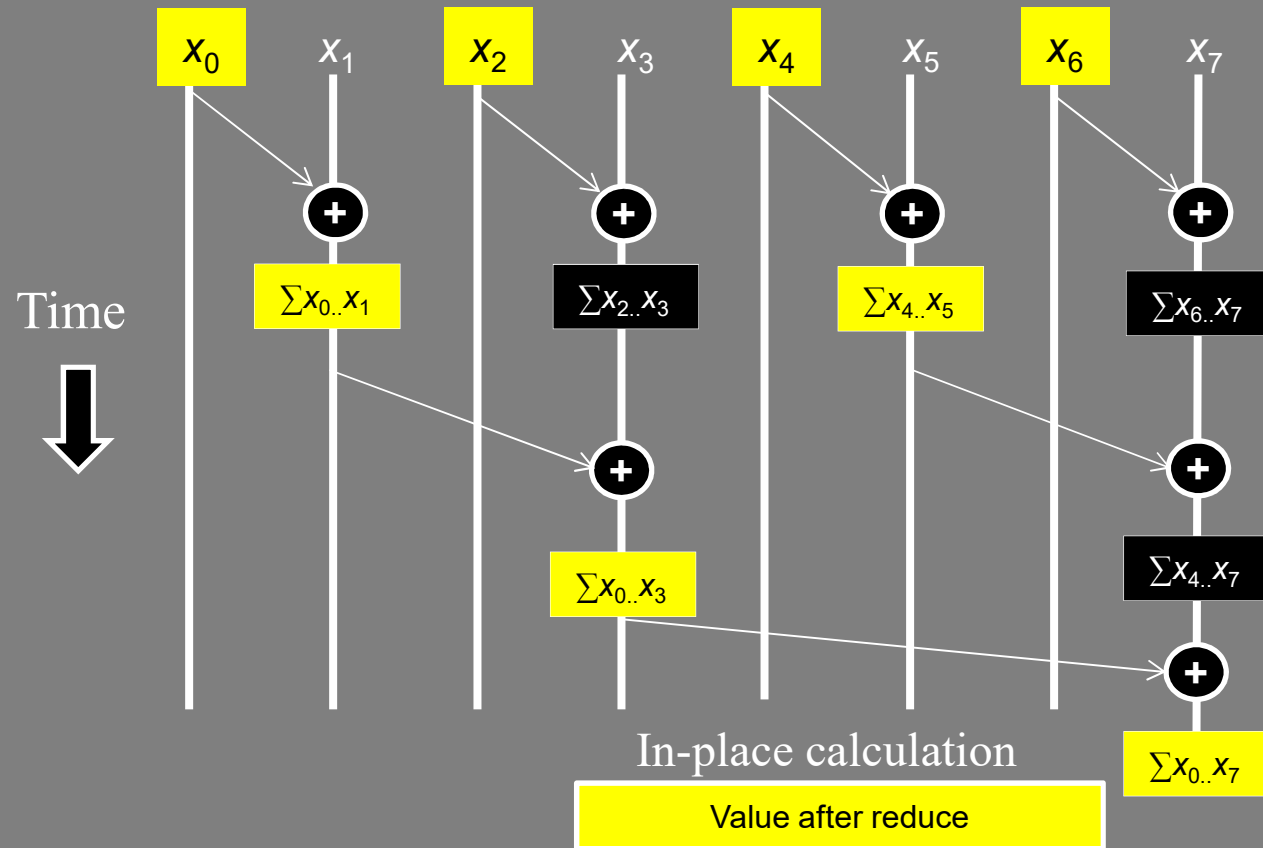
19

- This Scan executes $\log(n)$ parallel iterations
 - *The iterations do $(n-1)$, $(n-2)$, $(n-4)$, ..., $(n - n/2)$ adds each*
 - *Total adds: $n * \log(n) - (n-1) \rightarrow O(n * \log(n))$ work*
- This scan algorithm is not work efficient
 - *Sequential scan algorithm does n adds*
 - *A factor of $\log(n)$ can hurt: 10x for 1024 elements!*
- A parallel algorithm can be slower than a sequential one when execution resources are saturated from low work efficiency
- This means we can complete a scan in $O(\lg n)$ time if we have $O(n \lg n)$ processors

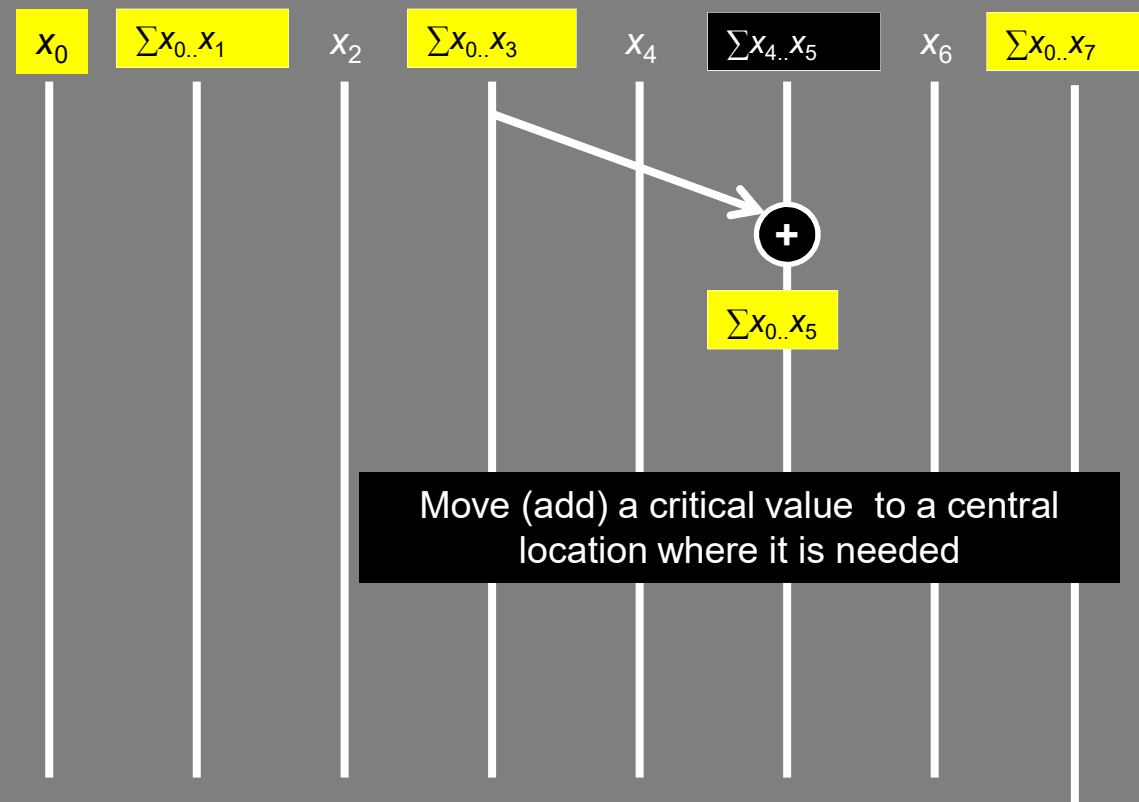


Parallel Scan - Reduction Phase

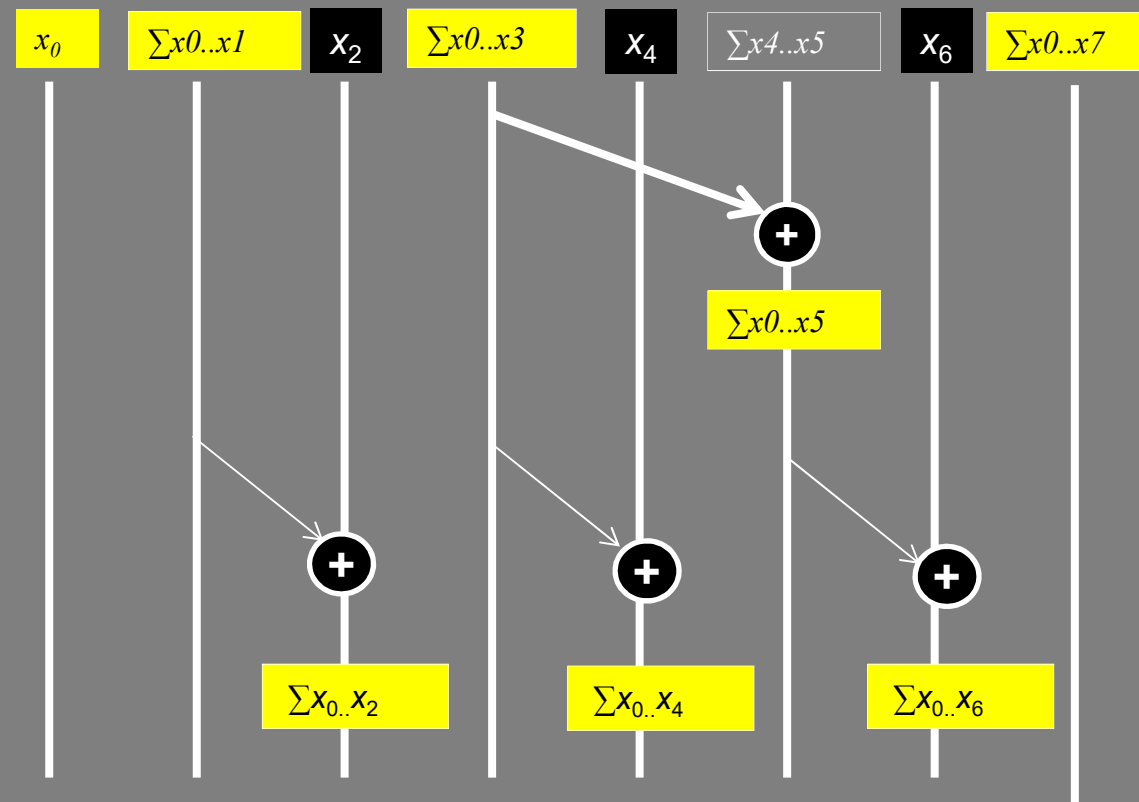
21



Parallel Scan - Post Reduction Reverse Phase




Parallel Scan - Post Reduction Reverse Phase²³



Final view

24



- Goal – build a scan algorithm that avoids the extra $O(\lg n)$ work
- Solution – make use of balanced binary trees
 - *Build a balanced binary tree on the input data and sweep the tree to compute the prefix sum*
- Consists of two parts
- Up Sweep
 - *Traverse tree from leaves to root computing partial sums during the traversal*
 - *The root node (last element in the array) will hold the sum of all elements*
- Down Sweep
 - *Traverse from root to leaves using the partial sums to compute the cumulative sums in place.*
 - *For the exclusive prefix sum, we replace the root with zero*

Blelloch Algorithm – Up Sweep

26

```
for(i = 0; i <= log_of_N_1; i++) {  
    two_i_p1 = 1 << (i+1);  
    two_i = 1 << i;  
  
    for(j = 0; j < N; j+=two_i_p1)  
        scanSum[j + two_i_p1 - 1] =  
            scanSum[j + two_i - 1] +  
            scanSum[j + two_i_p1 - 1];  
}
```

Blelloch Algorithm – Down Sweep

27

```
scanSum[N-1] = 0;
for(i = log_of_N_1; i >= 0; i--) {
    two_i_p1 = 1 << (i+1);
    two_i = 1 << i;

    for(j = 0; j < N; j+=two_i_p1) {
        long t = scanSum[j + two_i - 1];
        scanSum[j + two_i - 1] =
            scanSum[j + two_i_p1 - 1];
        scanSum[j + two_i_p1 - 1] = t +
            scanSum[j + two_i_p1 - 1];
    }
}
```

- This algorithm does complete the summation in $O(n)$ operations, giving the appearance of efficiency.
- Actually, the algorithm requires $2*(n-1)$ additions and $n-1$ swaps to run to completion.
- For large arrays, this algorithm will indeed out-perform the Hillis and Steele Algorithm.

a work-efficient scan kernel

29

- Two-phased balanced tree traversal
 - *Aggressive re-use of intermediate results*

A Scan Algorithm – Preview

30

3	1	7	0	4	1	6	3
---	---	---	---	---	---	---	---

Assignment : Parallelize using openMP.

Complexity of both the algorithms.

**Speedup for different size of the input array.
Also scalability?**

Problem

31

Old index

0	1	2	3	4	5	6	7	
0	1	0	0	1	1	0	1	
0	1	1	1	2	3	3	4	
1	4	5	7	6	3	2	0	↓

New index